

CO1 AT2

Student Name	:	Galla Sai Siddartha
Registration Number	:	192425093
Course Code & Name	:	MLA0302 - Reinforcement Learning
Date of Submission	:	21-07-2026

Question 1: Warehouse Robot Using Reinforcement Learning (MDP)

Aim :

To analyze how a warehouse robot can be modeled using a Markov Decision Process (MDP) and design a reward function that enables the robot to pick, pack, and transport items efficiently while ensuring speed, accuracy, and safety.

Markov Decision Process (MDP) Components:

1. States (S)

A state represents the current condition of the warehouse robot and its surroundings.

Examples:

Robot's current location.

2. Actions (A)

Actions are the possible operations the robot can perform.

Examples:

Move Forward

3. Rewards (R)

Rewards guide the robot toward the desired behavior.

Examples:

+20 for successfully picking an item.

4. Policy (π)

A policy is the strategy that tells the robot which action to take in every state.

Example:

If an item is nearby $\rightarrow$ Move toward it.

If carrying an item $\rightarrow$ Go to the packing station.

5. Environment

The environment is the warehouse where the robot operates.

It includes:

Warehouse layout.

Reward Function Design

A reward function balances **speed, accuracy, and safety**.

$R = 30 \times \text{Successful Delivery} + 20 \times \text{Correct Pick} + 5 \times \text{Fast Completion} - 10 \times \text{Collision} - 5 \times \text{Extra Steps} - 20 \times \text{Wrong Delivery}$

Reward Table

Event	Reward
Correct Item Picked	+20
Successful Delivery	+30
Fast Task Completion	+5
Efficient Route	+2
Collision with Obstacle	-10
Wrong Item Picked	-15
Wrong Delivery	-20

Influence of Reward Design

Positive Rewards

- Encourage quick task completion.
- Improve navigation efficiency.

Negative Rewards

- Prevent collisions.
- Reduce unnecessary movements.

Balanced Reward Design

A balanced reward function helps the robot:

- Learn the shortest safe path.
- Complete tasks faster.

Result

The warehouse robot is successfully modeled as a Markov Decision Process (MDP) using states, actions, rewards, policy, and environment. A well-designed reward function that balances speed, accuracy, and safety enables the robot to learn efficient navigation, avoid obstacles, minimize delays, and improve overall warehouse productivity.

Question 2 : Movie Recommendation System Using Reinforcement Learning

Aim :

To apply Reinforcement Learning (RL) concepts for a movie recommendation system by defining the state representation, action space, and reward function, and analyzing how reward signals and the exploration–exploitation trade-off affect user engagement.

Reinforcement Learning Components:

1. State (S)

A state represents the current information about the user.

Examples:

- User's viewing history.

2. Action (A)

Actions are the recommendations made by the system.

Examples:

- Recommend an Action movie.

3. Reward Function (R)

The reward measures how satisfied the user is with the recommendation.

Examples:

- +10 if the user watches the recommended movie.
- +20 if the movie is watched completely.

4. Policy (π)

The policy decides which movie should be recommended based on the current user state.

Example:

- If the user frequently watches action movies, recommend similar action movies.

5. Environment

The environment is the streaming platform where users interact with recommended movies.

It includes:

- Users.
- Movie catalog.

Reward Function Design

$R = 20 \times \text{Complete Watch} + 10 \times \text{Movie Click} + 5 \times \text{Positive Rating} + 8 \times \text{Continued Viewing} - 10 \times \text{Skip} - 15 \times \text{Exit Platform}$

Reward Table

Event	Reward
Movie Clicked	+10
Movie Watched Completely	+20
Positive Rating	+5
Continued Watching	+8
Recommendation Ignored	-5
Movie Skipped	-10
User Leaves Platform	-15

Effect of Different Reward Signals

Positive Rewards

- Encourage accurate recommendations.
- Increase watch time.

Negative Rewards

- Reduce poor recommendations.
- Prevent repeated suggestions of unwanted movies.

Balanced Reward Design

A balanced reward function helps the system:

- Recommend relevant movies.
- Increase user retention.

Advantages

- Personalized movie recommendations.
- Improves user engagement.
- Increases watch time.

Applications

- Netflix.
- Amazon Prime Video.
- Disney+.

Result :

The movie recommendation system is successfully modeled using Reinforcement Learning by defining appropriate states, actions, rewards, policy, and environment.

Question 3 : Autonomous Agricultural System Using Reinforcement Learning

Aim :

To design a Reinforcement Learning (RL) framework for an autonomous agricultural system that optimizes irrigation and fertilizer usage based on soil conditions.

Reinforcement Learning Components:

1. State (S)

A state represents the current condition of the farm.

Examples:

Soil moisture level.

2. Action (A)

Actions are the decisions taken by the RL agent.

Examples:

Increase irrigation.

3. Reward Function

The reward encourages efficient resource usage while maximizing crop growth.

Examples:

- **+20** for healthy crop growth.
- **+15** for optimal soil moisture.

4. Policy (π)

The policy determines the best action for each farm condition.

Example:

If soil moisture is low and no rain is expected → Start irrigation.

5. Environment

The environment is the agricultural field where the RL agent operates.

It includes:

- Soil conditions.

Reward Function Design

$R = 20 \times \text{Healthy Crop Growth} + 15 \times \text{Optimal Soil Moisture} + 10 \times \text{Efficient Water Usage} + 10 \times \text{Correct Fertilizer Use} - 10 \times \text{Overwatering} - 10 \times \text{Excess Fertilizer} - 20 \times \text{Crop Damage}$

Reward Table:

Event	Reward
Healthy Crop Growth	+20
Optimal Soil Moisture	+15
Efficient Water Usage	+10
Correct Fertilizer Application	+10
Overwatering	-10
Excess Fertilizer Usage	-10

Effect of Delayed Rewards

In agriculture, the results of irrigation and fertilizer decisions are usually seen after several days or weeks.

Positive Effects

- Helps maximize long-term crop yield.
- Encourages efficient use of water and fertilizer.

Challenges

- The agent does not receive immediate feedback.
- It becomes harder to identify which action caused a successful or poor harvest.

Advantages

- Optimizes water usage.
- Reduces fertilizer wastage.

Applications

- Smart irrigation systems.
- Precision farming.

Result

The autonomous agricultural system is successfully modeled using Reinforcement Learning by defining appropriate states, actions, rewards, policy, and environment. Although rewards are delayed until crop growth is observed, using a suitable algorithm such as Q-Learning (or Deep Q-Network for complex environments) enables the system to learn efficient irrigation and fertilizer strategies that maximize crop yield while minimizing resource consumption.

Question 4 : Cybersecurity Intrusion Detection Using Reinforcement Learning

Aim :

To design a Reinforcement Learning (RL) framework for a cybersecurity system that detects and responds to network attacks in real time by defining states, environment..

Reinforcement Learning Components

1. State (S)

A state represents the current condition of the network.

Examples:

- Network traffic volume.

2. Environment

The environment is the computer network where the RL agent operates.

It includes:

- Users and devices.

3. Actions (A)

Actions are the security responses chosen by the RL agent.

Examples:

- Allow network traffic.

4. Reward Function (R)

The reward function encourages accurate threat detection while minimizing false alarms.

Examples:

- +20 for correctly detecting an attack.

5. Policy (π)

The policy determines the best response for each network state.

Example:

- If abnormal traffic is detected → Block the IP address.
-

Reward Function Design

$R = 20 \times \text{Attack Detected} + 15 \times \text{Threat Blocked} + 10 \times \text{Data Protected} - 10 \times \text{False Alarm} - 15 \times \text{Missed Attack} - 20 \times \text{System Compromise}$

Reward Table

Event	Reward
Correct Attack Detection	+20
Successfully Blocked Attack	+15
Protected Sensitive Data	+10
Early Threat Detection	+5
False Alarm	-10

Challenges of Sparse Rewards

Sparse rewards occur because successful attacks or major security events happen infrequently.

Effects

The agent receives limited feedback.

Challenges of Real-Time Decision Making

The cybersecurity system must react immediately to threats.

Challenges

Large volumes of network traffic.

Advantages

- Detects attacks automatically.
- Learns new attack patterns over time.
- Improves network security.

Applications

- Intrusion Detection Systems (IDS).
- Intrusion Prevention Systems (IPS).
- Enterprise network security.

Result

The cybersecurity system is successfully modeled using Reinforcement Learning by defining suitable states, environment, actions, rewards, and policy. Although sparse rewards and real-time decision-making present challenges, a carefully designed reward function enables the RL agent to detect attacks efficiently, respond rapidly to threats, reduce false alarms, and continuously improve network security.

Experiment 5 : Ride-Sharing Driver Allocation and Dynamic Pricing Using Reinforcement Learning

Aim

To apply Reinforcement Learning (RL) for optimizing driver allocation and dynamic pricing in a ride-sharing system by defining states, actions, rewards, and environment, while analyzing the effects of environmental uncertainty and ethical concerns in pricing.

Reinforcement Learning Components

1. State (S)

A state represents the current condition of the ride-sharing system.

Examples:

- Number of available drivers.

2. Actions (A)

Actions are the decisions taken by the RL agent.

Examples:

- Assign the nearest driver.

3. Reward Function (R)

The reward encourages efficient driver allocation and fair pricing.

Examples:

- +20 for successful ride completion.

4. Policy (π)

The policy determines the best action for each system state.

Example:

- If demand is low $\rightarrow$ Reduce prices to attract customers.

5. Environment

The environment is the ride-sharing platform where the RL agent operates.

It includes:

- Drivers.

Reward Function Design

$R = 20 \times \text{Ride Completed} + 15 \times \text{Reduced Waiting Time} + 10 \times \text{Driver Utilization} + 10 \times \text{Customer Satisfaction} - 15 \times \text{Ride Cancellation} - 20 \times \text{Excessive Surge Pricing}$

Reward Table

Event	Reward
Successful Ride Completion	+20
Reduced Customer Waiting Time	+15
High Driver Utilization	+10
Customer Satisfaction	+10

Effect of Environmental Uncertainty

The ride-sharing environment changes continuously, making learning more challenging.

Sources of Uncertainty

- Sudden increase in ride requests.
- Traffic congestion.

Effects on Learning

- Demand patterns change frequently.
- Optimal pricing changes over time.

Ethical Concerns in Dynamic Pricing

Dynamic pricing should be fair and transparent.

Concerns

- Excessive surge pricing during emergencies.
- Higher fares for low-income users.

Solutions

- Set a maximum limit on surge pricing.
- Explain fare calculations clearly to users.

Advantages

- Reduces customer waiting time.
- Improves driver allocation.

Applications

- Uber.
- Ola.

Result

The ride-sharing system is successfully modeled using Reinforcement Learning by defining suitable states, actions, rewards, policy, and environment.

Question 6 : Smart Home Automation Using Reinforcement Learning

Aim

To design a Reinforcement Learning (RL) framework for a smart home automation system that controls lighting, heating, and cooling based on occupancy and weather conditions, while balancing user comfort and energy efficiency.

Reinforcement Learning Components

1. State (S)

A state represents the current condition of the smart home.

Examples:

- Number of people in the house.

2. Actions (A)

Actions are the decisions taken by the RL agent.

Examples:

- Turn lights ON.

3. Reward Function (R)

The reward encourages both user comfort and energy saving.

Examples:

- +20 for maintaining comfortable room temperature.

4. Policy (π)

The policy determines the best action based on the current home conditions.

Example:

- If a room is occupied $\rightarrow$ Turn lights ON.

5. Environment

The environment is the smart home where the RL agent operates.

It includes:

- Rooms.
-

Reward Function Design

$R = 20 \times \text{Comfort} + 15 \times \text{Energy Saving} + 10 \times \text{Proper Lighting} - 15 \times \text{User Discomfort} - 20 \times \text{Excess Energy Usage}$

Reward Table

Event	Reward
Comfortable Temperature	+20
Reduced Energy Consumption	+15
Proper Lighting Control	+10
Automatic Appliance Shutdown	+10
User Discomfort	-15
Lights ON in Empty Room	-10
Excess Energy Usage	-20
Unnecessary Heating/Cooling	-10

Handling Conflicting Objectives

The smart home must balance comfort and energy saving, which may sometimes conflict.

Comfort

- Maintain comfortable room temperature.

Energy Saving

- Reduce electricity usage.

Balancing Both Objectives

- Use a weighted reward function that gives importance to both comfort and energy efficiency.

Impact of Reward Shaping

Reward shaping provides additional intermediate rewards to guide learning.

Benefits

- Speeds up the learning process.
- Encourages energy-efficient behavior.

Advantages

- Reduces electricity consumption.
- Improves user comfort.

Applications

- Smart homes.
- Smart buildings.
- Office automation.

Question 7 : Strategy Game Playing Agent Using Reinforcement Learning

Aim

To model a strategy game as a Markov Decision Process (MDP) and design a reward structure that encourages long-term planning while analyzing the effect of exploration strategies on the RL agent's performance.

Markov Decision Process (MDP) Components

1. States (S)

A state represents the current situation of the game.

Examples:

- Position of all game units.

2. Actions (A)

Actions are the moves the RL agent can perform.

Examples:

- Move units.

3. Rewards (R)

Rewards encourage the agent to develop long-term winning strategies.

Examples:

- +50 for winning the game.

4. Policy (π)

The policy decides the best action in each game state.

Example:

- If enemy is weak $\rightarrow$ Attack.

5. Environment

The environment is the strategy game itself.

It includes:

- Game map.

Reward Function Design

$R = 50 \times \text{Game Won} + 20 \times \text{Enemy Defeated} + 15 \times \text{Objective Captured} + 10 \times \text{Resources Collected} - 20 \times \text{Poor Strategy} - 50 \times \text{Game Lost}$

Reward Table

Event	Reward
Win Game	+50
Defeat Enemy	+20
Capture Objective	+15
Collect Resources	+10
Lose Unit	-10
Poor Strategic Move	-20
Lose Game	-50

Encouraging Long-Term Strategy

The reward function should:

- Give higher rewards for winning the entire game rather than short-term actions.
- Encourage resource collection for future advantages.

Effect of Exploration Strategies

Exploration

- Tries new actions and strategies.
- Helps discover better game tactics.

Exploitation

- Uses the best-known strategy.
- Improves immediate performance.

Balanced Strategy

Using methods such as **ϵ -greedy exploration** allows the agent to:

- Explore new strategies with a small probability.

Advantages

- Learns complex game strategies.
- Improves decision-making over time.

Applications

- Chess AI.
- Go AI.

Result

The strategy game is successfully modeled as a Markov Decision Process (MDP) using states, actions, rewards, policy, and environment.

Question 8 : Fraud Detection Using Reinforcement Learning

Aim

To apply Reinforcement Learning (RL) for real-time fraud detection in financial transactions by defining the state space, action space, and reward function, while analyzing the challenges of imbalanced data.

Reinforcement Learning Components

1. State (S)

A state represents the current transaction details.

Examples:

- Transaction amount.

2. Action (A)

Actions are the decisions taken by the RL agent.

Examples:

- Approve transaction.

3. Reward Function (R)

The reward encourages accurate fraud detection while minimizing incorrect decisions.

Examples:

- +20 for correctly identifying fraudulent transactions.

4. Policy (π)

The policy decides the best action based on the current transaction state.

Example:

- If the transaction has a high fraud risk $\rightarrow$ Block or request OTP.

5. Environment

The environment is the banking or financial transaction system.

It includes:

- Customers.
- Bank accounts.

Reward Function Design

$R = 20 \times \text{Correct Fraud Detection} + 15 \times \text{Correct Approval} + 10 \times \text{Fraud Prevented} - 20 \times \text{Missed Fraud} - 15 \times \text{False Rejection} - 10 \times \text{Unnecessary Review}$

Reward Table

Event	Reward
Correct Fraud Detection	+20
Correct Genuine Approval	+15
Fraud Prevented	+10
Successful Authentication	+5
Missed Fraud	-20
Genuine Transaction Rejected	-15
Unnecessary Manual Review	-10
Financial Loss	-25

Challenges of Imbalanced Data

Fraudulent transactions are much fewer than genuine transactions.

Effects

- The RL agent receives fewer fraud examples.
- Learning fraud patterns becomes difficult.

Challenges of Delayed Feedback

Fraud is often confirmed only after customer complaints or investigations.

Effects

- Rewards are delayed.
- The agent takes longer to learn correct decisions.

False Positive (Genuine Transaction Blocked)

- Customer inconvenience.
- Loss of customer trust.

False Negative (Fraud Not Detected)

- Financial loss.
- Data theft.

Advantages

- Detects fraud in real time.
- Learns new fraud patterns automatically.

Applications

- Online banking.
- Credit card fraud detection.
- Digital payment systems.

Question 9 : Personalized Rehabilitation System Using Reinforcement Learning

Aim

To design a Reinforcement Learning (RL) framework for recommending personalized rehabilitation exercises based on patient progress by defining states, actions, rewards, and environment, while analyzing delayed rewards, patient variability, and ethical and safety considerations.

Reinforcement Learning Components

1. State (S)

A state represents the current health condition of the patient.

Examples:

- Patient age.

2. Actions (A)

Actions are the rehabilitation exercises recommended by the RL agent.

Examples:

- Recommend stretching exercises.

3. Reward Function (R)

The reward encourages safe recovery and patient improvement.

Examples:

- +20 for improvement in recovery.

4. Policy (π)

The policy determines the best exercise recommendation for each patient state.

Example:

- If recovery is improving → Increase exercise intensity slightly.

5. Environment

The environment is the healthcare rehabilitation system.

It includes:

- Patients.

Reward Function Design

$R = 20 \times \text{Recovery Improvement} + 15 \times \text{Exercise Completed} + 10 \times \text{Pain Reduction} + 10 \times \text{Mobility Improvement} - 15 \times \text{Health Deterioration} - 20 \times \text{Unsafe Recommendation}$

Reward Table

Event	Reward
Recovery Improvement	+20
Exercise Completed Successfully	+15
Reduced Pain	+10
Improved Mobility	+10
Patient Discomfort	-10
Health Deterioration	-15
Unsafe Exercise Recommendation	-20
Missed Therapy Session	-5

Impact of Delayed Rewards

Recovery results are usually seen after several days or weeks.

Effects

- The RL agent receives feedback slowly.
- Learning takes more time.

Impact of Patient Variability

Each patient responds differently to rehabilitation.

Challenges

- Different ages and fitness levels.
- Different injury types.

Ethical Concerns and Safety Considerations

Ethical Concerns

- Protect patient privacy and medical data.
- Ensure fairness in treatment recommendations.

Safety Considerations

- Prevent exercises that may worsen injuries.
- Monitor patient pain and vital signs continuously.

Advantages

- Provides personalized rehabilitation plans.
- Improves patient recovery.

Applications

- Physical rehabilitation.
- Stroke recovery.

Question 10 : Smart Grid Electricity Management Using Reinforcement Learning

Aim

To analyze how a smart grid can be modeled as a Markov Decision Process (MDP) for efficient electricity distribution and peak load reduction, and to design a reward function that balances efficiency and reliability.

Markov Decision Process (MDP) Components

1. States (S)

A state represents the current condition of the smart grid.

Examples:

- Current electricity demand.

2. Actions (A)

Actions are the decisions made by the RL agent.

Examples:

- Increase power supply.

3. Reward Function (R)

The reward encourages efficient and reliable electricity distribution.

Examples:

- +20 for meeting electricity demand.

4. Policy (π)

The policy determines the best action for each grid state.

Example:

- If demand is high $\rightarrow$ Supply more electricity.

5. Environment

The environment is the smart electricity grid.

It includes:

- Power generation units.

Reward Function Design

$R = 20 \times \text{Demand Met} + 15 \times \text{Peak Load Reduced} + 10 \times \text{Efficient Energy Use} + 10 \times \text{Renewable Energy Usage} - 15 \times \text{Power Outage} - 20 \times \text{Grid Overload}$

Reward Table

Event	Reward
Electricity Demand Met	+20
Peak Load Reduced	+15
Efficient Energy Usage	+10
Renewable Energy Utilized	+10
Energy Wastage	-10
Power Outage	-15
Grid Overload	-20
Voltage/Frequency Instability	-10

Balancing Efficiency and Reliability

Efficiency

- Reduce energy wastage.
- Optimize electricity distribution.

Reliability

- Ensure uninterrupted power supply.
- Prevent overloads.

Challenges of Scalability

As the smart grid expands, the RL system becomes more complex.

Challenges

- Large number of consumers.
- Multiple power generation sources.

Challenges of Real-Time Learning

The RL agent must make decisions quickly as electricity demand changes.

Challenges

- Rapid fluctuations in power demand.
- Changing weather affecting renewable energy.

Advantages

- Optimizes electricity distribution.
- Reduces peak load demand.

Applications

- Smart power grids.
- Renewable energy management.